



# Concurrent Object-Oriented Programming (Part One)



Concurrent and Parallel Programming

# Shared memory and variables

---

- ▶ We know that **shared and mutable** memory regions represent dangers in multi-threaded programs. Synchronization tools have to be used for correct access.
- ▶ But, how this consideration **impact on the source code** of an object oriented program?
- ▶ At the code level, **which are the types of variables that can be shared between threads** and what kind of risk are associated?

# Variables and shared memory

---


- ▶ **Local variables and parameters:** no problem. Are on the stack.
- ▶ **Final static variables:** no problem. Are constants.
- ▶ **Static (non final) variables:** are global variables at the class level, always at risk of being shared. Are considered the nightmare of concurrent programming. Should be avoided.
- ▶ **Primitive type fields:** are shared when inside shared objects. Might be problematic.
- ▶ **Reference type fields:** behave like primitive type fields. When shared, the referenced array/object is also shared. Warning: an array/object can also be shared by copying the reference.

# State of an object

---

- ▶ **The state of an object** is composed of the values of its fields:
  - ▶ If they are all primitive type fields, their values represent the entire state of the object.
  - ▶ If some fields are references to other objects, the state may include the values of the fields of the referenced object too.


# State space - review



Every object has its own **state-space**. It's the space of all the **consistent states** in which the object can be.

The state-space is delimited by **the invariants**: rules that have to be respected by the object to avoid falling in an **inconsistent state**.

# Invariants - review



Invariant

example:

$0 \leq x < 1000 \ \&\&$

$0 \leq y < 10 \ \&\&$

$x * y \neq 5000$

Invariants represent the **limits of the values** that can be assumed by the variables (both static and instance).

Invariants have to be established by the constructors and have to be preserved by the methods.

# Example

```
public class NumberRange {  
    // INVARIANT: lower <= upper  
    private volatile int lower = 0;  
    private volatile int upper = 0;  
  
    public void setLower(int i) {  
        if (i > upper)  
            throw new IllegalArgumentException("can't set lower to " + i + " > upper");  
        lower = i;  
    }  
  
    public void setUpper(int i) {  
        if (i < lower)  
            throw new IllegalArgumentException("can't set upper to " + i + " < lower");  
        upper = i;  
    }  
  
    public boolean isInRange(int i) {  
        return (i >= lower && i <= upper);  
    }  
}
```

# State-space and invariants

---

- ▶ There are invariants that constrain **more than one** state **variable** at the same time. For example, it is the case of dependent variables in **compound actions**.
- ▶ The smaller the state-space, the easier it is to manage the state of the object.
- ▶ As we will see later, it is possible to use the **“final”** keyword to reduce the state-space of an object.



# Example



```
public class NumberRange {
    // INVARIANT: lower <= upper
    private volatile int lower = 0;
    private volatile int upper = 0;

    public void setLower(int i) {
        // Warning -- unsafe check-then-act
        if (i > upper)
            throw new IllegalArgumentException("can't set lower to " + i + " > upper");
        lower = i;
    }

    public void setUpper(int i) {
        // Warning -- unsafe check-then-act
        if (i < lower)
            throw new IllegalArgumentException("can't set upper to " + i + " < lower");
        upper = i;
    }

    public boolean isInRange(int i) {
        return (i >= lower && i <= upper);
    }
}
```

# Example

```
public class NumberRange {  
    // INVARIANT: lower <= upper  
    private volatile int lower = 0;  
    private volatile int upper = 0;  
  
    public void setLower(int i) {  
        // Warning -- unsafe check-then-act  
        if (i > upper)   
            throw new IllegalArgumentException("can't set lower to " + i + " > upper");  
        lower = i;  
    }  
  
    public void setUpper(int i) {  
        // Warning -- unsafe check-then-act  
        if (i < lower)   
            throw new IllegalArgumentException("can't set upper to " + i + " < lower");  
        upper = i;  
    }  
  
    public boolean isInRange(int i) {  
        return (i >= lower && i <= upper);  
    }  
}
```

There's not  
enough  
protection!

# Concurrency and objects

---

- ▶ In your opinion, when and how should the state of an object be protected in **a multi-threaded application?**
- ▶ Which approach do you propose to use?

# Thread-safe classes and objects

---

- ▶ One of the most frequently used approaches is to implement and use **thread-safe classes and objects**.
- ▶ An object is said to be thread-safe if it **can be accessed from multiple threads at the same time**:
  - ▶ independently from the execution order (interleaving) of the operations executed by the threads.
  - ▶ **without using synchronization tools (e.g. locks) in the client code that uses the object.**

# Stateless objects

---

- ▶ **Stateless objects** are always thread-safe.
- ▶ Possible state exists only in parameters and local variables of the methods, which live in the thread stack.

# Example (1 / 2)

```
class StatelessObject {
    public double doTheWork(final double r) {
        final double volume = (4.0 / 3.0) * Math.PI * r * r * r;
        System.out.println("Volume for r: " + r + " is = " + volume);
        return volume;
    }
}

class Runner implements Runnable {
    private final StatelessObject statelessObject;
    private final int r;

    public Runner(final StatelessObject statelessObject, final int r) {
        this.statelessObject = statelessObject;
        this.r = r;
    }

    @Override
    public void run() {
        statelessObject.doTheWork(r);
    }
}
```

# Example (2/2)

```
public class TestStatelessObject {
    static final ArrayList<Thread> threads = new ArrayList<Thread>();
    static final StatelessObject statelessObject = new StatelessObject();

    public static void main(final String[] args) {
        for (int i = 0; i < 10; i++)
            threads.add(new Thread(new Runner(statelessObject, i)));

        try {
            for (final Thread t : threads)
                t.start();
            for (final Thread t : threads)
                t.join();
        } catch (final InterruptedException e) {
            return;
        }
    }
}
```

# Statefull objects

- ▶ For statefull objects, **thread-safety** can be obtained in the following 3 ways:
  - ▶ By transforming them into **immutable** objects.  
**Immutable objects** are always **thread-safe**.
  - ▶ **By avoiding to share the** object state, i.e. by applying **confinement** techniques.
  - ▶ By using **synchronization tools** (volatile variables, atomic variables, locks, ...) inside the object.



# Non-thread-safe objects

---

- ▶ A thread-safe program does NOT mandatorily require to be composed of just thread-safe objects.
- ▶ A thread-safe program can also contain non-thread-safe objects.
- ▶ For example, it might be the case of objects always accessed from just one of the threads; or objects used with synchronization (e.g. locks) in the calling client code.

# Development of thread-safe classes

---

- ▶ Encapsulation and data hiding techniques are very helpful, they allow to **isolate and hide** the state variables of the object.
- ▶ **Encapsulation** helps assessing whether a class is thread-safe or not, without analyzing the whole program.

# Development of thread-safe classes

---

- ▶ The process of designing a **thread-safe class** includes:
  - ▶ Identification of the variables that represent the **state of the object**.
  - ▶ Identification of the **invariants** that limit the state-space.
  - ▶ Definition of rules called **synchronization policies** for correctly managing concurrent accesses to the state variables.

# Encapsulation and synchronization

---

- ▶ The constraints imposed by the invariants generate **synchronization requirements**. Operations that may cause invalid transitions or other types of errors must be made **atomic**.
- ▶ On the other hand, if there are no constraints, encapsulation and synchronization requirements can be relaxed to obtain **greater flexibility and better performances**.

# Synchronization policy

---

- ▶ **The synchronization policy** defines how an object coordinates the multi-threaded access to its state, by a combination of:
  - ▶ **Immutability**
  - ▶ **Confinement**
  - ▶ **Locking and other synchronization tools**
- ▶ Developing a thread-safe class means ensuring that its **invariants are not violated** under concurrent accesses.

# Read-only shared state

---

- ▶ In case of shared state that is **read-only**, there's no need for synchronization.
- ▶ Values **are never changed**, therefore there's no risk for visibility problems (registers/caches with inconsistent content) or race-conditions.
- ▶ In multi-threaded object-oriented programs, it is possible to take advantage of this feature!
- ▶ The Java language has specific support for read-only state variables by means of the **"final" keyword**.

# Final fields - review

---

- ▶ A **final field** can only be initialized once, either by **default initialization** or by **an assignment in the constructor (or static initializer for static fields)**.
- ▶ Each final field must be initialized at the latest before the end of the constructor or static initializer of the class in which is declared.
- ▶ **ATTENTION:** for reference types, it's the reference itself that is final. Not the referenced object!

# Initialization safety

---

- ▶ The Java memory model provides specific support for the **initialization safety of final fields**.
- ▶ Without initialization safety, there could be the risk of accessing objects still **under construction**.
- ▶ Thanks to initialization safety, it is possible to access shared (read-only) fields even without synchronization.
- ▶ **ATTENTION:** if a final field is a reference to a mutable object, synchronization is then required (when accessing the mutable object).



# Immutable objects

- ▶ If the mutability of an object is fully avoided, the result is an **immutable object**.
- ▶ An object is **immutable** if it is not possible to modify its state after the object has been constructed.
- ▶ **Final fields are used to** implement immutable objects. Thanks to **initialization safety**, it is possible to access shared immutable objects without synchronization.
- ▶ However, to develop **immutable** objects it is **not sufficient to just set all fields to final** because of the potential presence of final references to mutable objects.

# Immutable objects

---

To obtain an immutable object:

- ▶ Specify all fields as **private and final**.
- ▶ Don't let subclasses override methods. The easiest way is to declare **the class or its methods as final**. A more sophisticated approach is to declare the constructor as private and manage the creation of objects with factory methods.
- ▶ If the fields contain references to mutable objects, prevent these objects from being modified:
  - ▶ No code in the methods that modify the contained mutable objects.
  - ▶ Never share references to mutable objects. Always make protection copies. Use the copy-on-write approach.

# Immutable objects: example

---

```
public final class Planet {  
  
    /** Final primitive data is always immutable. */  
    private final double fMass;  
  
    /** An immutable object field. */  
    private final String fName;  
  
    /** A mutable object field. */  
    private final Date fDateOfDiscovery;  
  
    public Planet(final double aMass, final String aName, final long discoveryTime) {  
        fMass = aMass;  
        fName = aName;  
        fDateOfDiscovery = new Date(discoveryTime);  
    }  
  
    ...  
}
```

# Immutable objects: example

...

```
public double getMass() {  
    return fMass;  
}
```

```
// String is immutable and cannot be changed.  
public String getName() {  
    return fName;  
}
```

```
// Returns a defensive copy of the field.  
public Date getDateOfDiscovery() {  
    return new Date(fDateOfDiscovery.getTime());  
}
```

```
public Planet rediscoverPlanet(final long newDiscoveryTime) {  
    return new Planet(fMass, fName, newDiscoveryTime);  
}
```



copy-on-write

# Confinement techniques

---

- ▶ In addition to the use of synchronization tools, several **specific techniques** can be exploited for the development of **multi-threaded applications**.
- ▶ These techniques are very effective and should always be taken into consideration when developing a concurrent program.
- ▶ Many of these techniques have a strong impact on the architectural design of the application. It is advised to apply them already in early design stages to avoid the need of significant refactoring.

# Confinement techniques

---

- ▶ As discussed, in a multi-threaded application it is required to correctly manage **the state that is shared and mutable**.
- ▶ But if there is state that is not shared and mutable, no protection is needed!
- ▶ Consequently, **if it is possible to avoid sharing variables and fields**, things are easier.
- ▶ For this reason the following techniques can be used:
  - ▶ Thread Confinement
  - ▶ Stack Confinement
  - ▶ Instance Confinement

# Thread Confinement

- ▶ **Thread confinement** is one of the easiest ways to achieve thread-safety.
- ▶ If a field or referenced object is **confined to the use of a single thread**, the execution from the point of view of that field/object can be compared to the execution of a mono-threaded application, ... therefore, no risk of concurrency problems has to be considered.

# Thread Confinement

---

- ▶ During the confinement time, it's important to ensure that the objects have no chance to escape (undesired sharing of objects).
- ▶ For the confined objects, it is important to ensure that they are also **created** by the same thread.
- ▶ The confinement might also be **temporary**. Later, it could be interrupted by **sharing the values or the objects with a synchronization tool** (e.g. a lock).
- ▶ Thread confinement is a design element of the program. To obtain it, **there's NO language-specific solution**. It must be implemented with discipline.



# Example

```
class Counter implements Runnable {
    private int value = 0;

    @Override
    public void run() {
        while (true) {
            value++;
            if (value % 10000 == 0) System.out.println(value);
            if (value == 100000) return;
        }
    }
}

public class CounterExample {
    static final ArrayList<Thread> threads = new ArrayList<>();

    public static void main(final String[] args) {
        for (int i = 0; i < 10; i++)
            threads.add(new Thread(new Counter()));
        for (final Thread t : threads)
            t.start();
        try {
            for (final Thread t : threads)
                t.join();
        } catch (final InterruptedException e) { return; }
    }
}
```

# Stack Confinement

- ▶ **Stack confinement** is a special version of thread confinement, where a value or referenced object can **only** be accessed **through method parameters or local variables**.
- ▶ Parameters and local variables exist on the thread stack, which is not accessible from other threads, therefore are **automatically confined to the running thread**.
- ▶ Stack confinement is easier and more intuitive to implement than thread confinement and results in less fragile code.
- ▶ Parameters and local variables of **primitive type are always stack confined**. For **reference types**, it is required to **ensure that the reference never escapes the thread**.

```
public class StackConfinementExample {  
    public int computeOperation(final int value1, final int value2) {  
        int y = value1;  
        for (int i = 0; i < 10; i++) {  
            y *= value2;  
        }  
  
        if (y < 0)  
            return 0;  
        return y;  
    }  
}
```


# Instance confinement

---

- ▶ One of the most effective techniques is to **confine the values and referenced objects** that compose the state of an object **directly in its instance**.
- ▶ Thread-safety can then be obtained by ensuring that:
  - ▶ the instance (container object) is accessed only by a single thread (thread confinement), or
  - ▶ all accesses to the instance are protected by synchronization means (i.e. it is a thread-safe object).
- ▶ **Encapsulation** is of strong help to achieve instance confinement. All code paths that access an encapsulated value or object are fully under control of the class.

# Example

```
public class PersonSet {  
    private final Set<Person> mySet = new HashSet<>();  
  
    public synchronized void addPerson(Person p) {  
        mySet.add(p);  
    }  
  
    public synchronized boolean containsPerson(Person p) {  
        return mySet.contains(p);  
    }  
  
    interface Person {  
        // ...  
    }  
}
```



# Pattern for thread-safe objects

- ▶ The concept of instance confinement introduces a possible design pattern:

*encapsulate all mutable states in an object and protect them from concurrent accesses by using synchronized methods:*

- ▶ all the mutable state of the object is encapsulated
  - ▶ all the mutable state of the object is protected by the intrinsic lock of the object
- ▶ This design pattern is simple but often not very effective in terms of performance.
- ▶ For **better performances** other types of locks (such as RW-locks), or for example atomic variables, might be used.

# Example

```
public class SynchronizationPatternExample {  
    private String name = null;  
    private boolean nameModified = false;  
  
    public synchronized String getName() {  
        return name;  
    }  
  
    public synchronized boolean isNameModified() {  
        return nameModified;  
    }  
  
    public synchronized void unsetNameModified() {  
        this.nameModified = false;  
    }  
  
    public synchronized void setName(final String name) {  
        this.name = name;  
        nameModified = true;  
    }  
}
```

# One writer - many readers

- ▶ Another design pattern that profits from the confinement techniques, but allows data to be shared between threads is called **one writer - many readers**.
- ▶ In this solution, just a **single thread is allowed to modify (write) the shared state**, while other threads have **only rights to read it**. The consequence is that no read-modify-write race-condition might ever happen.
- ▶ This pattern, can be implemented with just one **volatile variables** for sharing the state. Keep in mind that the visibility effect of volatile variables also extends to the other shared variables (if things are done correctly).




# One writer - many readers

---

- ▶ To increase the effectiveness of this approach, it is recommended that the only thread that modifies the state uses **thread/stack confinement until the instant of sharing**.
- ▶ For the threads that read the state, to avoid check-then-act race-conditions, it is also recommended to **switch to thread/stack confinement immediately after the initial read of the shared state**.

# Example

It has to be ensure that  
only one thread is  
allowed to invoke this  
method!



```
public class OneWriterExample {  
    private volatile int x = 1;  
  
    public void modifyValue(final int value) {  
        int y = x;  
        y *= value;  
        if (y < 0)  
            return;  
        x = y;  
    }  
  
    public int readValue() {  
        int tempX = x;  
        if (tempX > 1000)  
            return 1000;  
        else  
            return tempX;  
    }  
}
```

# Summary of topics

---

- ▶ Shared memory and types of variables
- ▶ Objects: state-space and invariants
- ▶ Thread-safe objects and classes
- ▶ The added value of encapsulation
- ▶ Final fields and immutable objects
- ▶ Confinement techniques
- ▶ Design pattern for thread-safe objects
- ▶ The one-writer many-readers pattern