



Concurrent Object-Oriented Programming (Part Two)



Concurrent and Parallel Programming

Construction, publication and sharing

- ▶ **Correctly encapsulating and managing the state of** objects, by possibly making them thread-safe, is not sufficient to create multi-threaded object oriented programs.
- ▶ For correct concurrent OOP, robust techniques are also required to:
 - ▶ **construct the objects,**
 - ▶ **make the objects accessible (publication),**
 - ▶ **share the use of the objects.**

Safe construction

- ▶ If a **reference to an object escapes** while the object is under construction, it may happen that the object gets shared between the threads, **when it is still not fully and correctly constructed (therefore in an inconsistent state)**.
- ▶ In particular, care has to be taken **to avoid the escape of the "this" reference** during construction time. For example, it might happen if a thread is started by the constructor code.

Example

```
public class EventListener {  
    public EventListener(EventSource eventSource) {  
        // do our initialization  
        // ...  
  
        // register ourselves with the event source  
        eventSource.registerListener(this);  
    }  
  
    public void onEvent(Event e) {  
        // handle the event  
    }  
}
```

The EventSource is managed by
an independent thread

```
public class RecordingEventListener extends EventListener {  
    private final List<Event> list;  
  
    public RecordingEventListener(EventSource eventSource) {  
        super(eventSource);  
        list = Collections.synchronizedList(new ArrayList<>());  
    }  
  
    public void onEvent(Event e) {  
        list.add(e);  
        super.onEvent(e);  
    }  
  
    public Event[] getEvents() {  
        return list.toArray(new Event[list.size()]);  
    }  
}
```

Safe construction

- ▶ Obtaining the safe construction of a class, **might be difficult because of inheritance**. To avoid mistakes it has to be avoided that extensions of the class unintentionally violate the safe construction constraint.
- ▶ One of the possible solutions is to limit inheritance, by specifying the class as final, or by using a **private constructor** and a **public factory method**.

Example

```
class AccumulatorWithFM {
    static int initialValue = 0;
    private int value = 0;

    static AccumulatorWithFM createAccumulator(final int value) {
        if (value > initialValue)
            return new AccumulatorWithFM(value);
        else if (value > 0)
            return new AccumulatorWithFM(initialValue);
        else
            return null;
    }

    private AccumulatorWithFM(final int value) {
        this.value = value;
    }

    // ...
}

class AccumulatorExtension extends AccumulatorWithFM {
    public AccumulatorExtension(final int value) {
        super(value); ← Error: the constructor is not visible!
    }
}
```

Publication of objects

- ▶ In multi-threaded applications it is important to analyze the state that a class **owns** (state ownership), in particular in terms of referenced objects.
- ▶ **State ownership** is an element of the class design. If a class references an *HashMap*, the state might include the *Map*, all *Map.Entry* instances, and potentially all the keys and objects contained in the *Map*.

Publication of objects

- ▶ **Publishing** an object means making it accessible to the code outside the scope of the class (or the method, in case of local variables), that owns it in a specific moment.
- ▶ **Attention:** the publication of an object is performed by a thread that makes the object available to other threads.

Example

```
class Printer implements Runnable {
    @Override
    public void run() {
        while (PubblicationExample.list == null) {
        }
        for (final String string : PubblicationExample.list) {
            System.out.print(string);
        }
    }
}

public class PubblicationExample {
    static List<String> list = null;


    public static void main(final String[] args) {
        // Starts thread here ...
        List<String> tempList = new ArrayList<String>();
        tempList.add("Goodmorning");
        tempList.add("students,");
        tempList.add("welcome ");
        tempList.add("to the ");
        tempList.add("course.");
        list = tempList;
    }
}
```

Publication of objects

- ▶ Publishing the internal state of a class can **compromise its encapsulation** and make it difficult to **preserve its invariants**.
- ▶ An object that is wrongly published (when it shouldn't), is called an **escaped object**.
- ▶ Any published object might be **used inappropriately by other threads**. For this reason it is important to limit the publication of objects, by profiting from **confinement techniques** as much as possible.

Confinement and publication

- ▶ If an object is confined to a context (e.g. a class by means of instance confinement, or a method by means of stack confinement), publishing and allowing the object to escape has to be considered a **BUG!**



Indirect and hidden publication

- ▶ The publication of an object might result in the **indirect publication of other objects** (objects that can be reached by following a chain of references).
- ▶ The publication of objects can be **hidden**.
- ▶ From the point of view of a class, **methods with a behavior that is not fully specified by the class** (such as overridden methods) are a danger.
- ▶ If a confined object is provided to an **overridden method**, **it has to be considered as being published**.


Safe publication

- ▶ *For the moment let's focus on the publication, not on any subsequent share and modification of the published objects!*
- ▶ **Safe publication** of an object means that the reference to the object and its state are **made accessible simultaneously (in an atomic way)** => this can only be done by using synchronization tools.
Visibility problems must be avoided!
- ▶ **All mutable objects must be safely published.**
Synchronization is required on both sides: the publisher, as well as the recipient of the publication.

Safe publication

Thread 1

Thread 2



Safe publication

- ▶ A safely constructed object **can be safely published**:
 - ▶ by initializing a static reference **with default initialization or a static initializer** (*public static MyObject obj = new MyObject(42);*)
 - ▶ by storing the object reference in a **volatile variable** (or in an **AtomicReference**)
 - ▶ by storing the object reference in a **final field** of a safely constructed object
 - ▶ by storing the object reference in a **lock-protected field**
 - ▶ by storing the object reference in a **thread-safe collection** (will be discussed later)

Example

Stack confinement to
avoid successive
inconsistencies

```
class Printer implements Runnable {  
    @Override  
    public void run() {  
        while (PubblicationExample.list == null) {  
        }  
        List<String> tempList = PubblicationExample.list;  
        for (final String string : tempList) {  
            System.out.print(string);  
        }  
    }  
}
```

Safe publication

```
public class PubblicationExample {  
    volatile static List<String> list = null;  
  
    public static void main(final String[] args) {  
        // Starts thread here ...  
        List<String> tempList = new ArrayList<String>();  
        tempList.add("Hello");  
        tempList.add("students, ");  
        tempList.add("welcome ");  
        tempList.add("«to ");  
        tempList.add("the course.");  
        list = tempList;  
    }  
}
```


Effectively immutable objects

- ▶ A safely published object can be accessed **in a read-only form without synchronization**, even if the object is not immutable. In this case, we are talking about **effectively immutable objects**.
- ▶ Why? Remember: **the visibility effect** of volatile variables, atomic variables, and locks, **also extends to any other state (including the internal state of published objects)** after using the volatile/atomic variable or the lock.
- ▶ At a practical level, it is possible to take advantage from effectively immutable objects, for example, in combination with the "one writer - many readers" technique.

Publication and mutable objects

- ▶ For all objects that are modified after their publication, safe publication **only** guarantees **correct visibility of the published state**.
- ▶ After publication, **mutable objects** have then to always be accessed safely. Therefore, either they are **thread-safe objects** or all accesses have to be protected by locks in the client code.

Publication and thread-safe objects

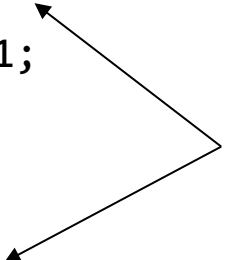
- ▶ If an object owned by a class is thread safe, does not participate in any invariant (e.g. dependencies with other variables), then it **can be published (safely) without difficulties**.
- ▶ Otherwise, ... **care has to be taken by introducing additional synchronization**, to avoid later inappropriate use of the object.

- ▶ Attention: even if an object is thread-safe, it might still be needed to use synchronization in the client code.

```
public static Object getLast(final List list) {  
    synchronized(list) {  
        final int lastIndex = list.size() - 1;  
        return list.get(lastIndex);  
    }  
}
```

```
public static void deleteLast(final List list) {  
    synchronized(list) {  
        final int lastIndex = list.size() - 1;  
        list.remove(lastIndex);  
    }  
}
```

A thread-safe
List is used



Possible publications of objects

- ▶ To summarize:
 - ▶ **Immutable objects:** can be published by any means, but it is better to publish them safely (for correct visibility of the reference).
 - ▶ **Effectively immutable objects:** must be published safely.
 - ▶ **Mutable objects:** must be published safely and must be thread-safe or protected by synchronization when used.

Possible object shares

► To summarize:

- **Thread-confined objects:** are accessed (read and write) only by the owning thread.
- **Shared read-only objects:** can be read concurrently without synchronization. Include immutable objects as well as effectively immutable objects.
- **Shared thread-safe objects:** perform synchronization internally. Can be accessed (read and write) concurrently without synchronization in the client code.
- **Shared objects protected by locks:** can only be accessed (read and write) when holding the associated lock.

Immutable Holder of Objects

- ▶ You might have the impression that immutable objects limit the possibility of updating the program states.
- ▶ But this is not true.
- ▶ Immutable objects can be **replaced with new instances containing the new values**, when needed.

Immutable Holder of Objects

- ▶ Immutable objects can be exploited to provide **atomicity**: every time a group of values/objects has to be used in an atomic way, an **immutable holder** of these values/objects can be created.
- ▶ Then, if the values/objects need to be modified, a **new holder** can be created to **replace the previous one**. This operation has to be the direct responsibility of the holder (**copy-on-write**).
- ▶ All threads that still have a reference to the old holder, continue to see a consistent (but old) state.

Immutable Holder of Objects

- ▶ To ensure correct visibility, this approach can be complemented with a **volatile (or atomic) reference to the holder instance** for performing safe-publication each time the holder instance is replaced with a new one.
- ▶ **Attention:** any race-condition such as check-then-act or read-modify-write has to be avoided.
- ▶ An effective solution to share the immutable holder is to use the **“one writer - many readers” technique**.

Example

```
//Immutable Holder
final class Holder {
    private final int n;
    private final int m;

    public Holder(final int n, final int m) {
        this.n = n; this.m = m;
    }

    public Holder incrementValues(final int delta) {
        return new Holder(n + delta, m + delta);
    }

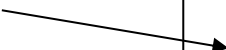
    //...
}
```

```
class Foo {
    private volatile Holder holder;

    public Holder getHolder() {
        return holder;
    }

    public void incrementValues(final int delta) {
        holder = holder.incrementValue(delta);
    }
}
```

In this case the substitution
is not atomic.



```
final class Foo {  
    private final AtomicReference<Holder> holderRef  
        = new AtomicReference<Holder>();  
  
    public Holder getHolder() {  
        return holderRef.get();  
    }  
  
    public void incrementValue(final int delta) {  
        holderRef.set(holderRef.get().incrementValue(delta));  
    }  
}
```



Also in this case the substitution is not atomic. It could eventually be made atomic by implementing the CAS idiom.

```
final class Foo {  
    private Holder holder;  
  
    public synchronized Holder getHolder() {  
        return holder;  
    }  
  
    public synchronized void setValue(final int n) {  
        holder = new Holder(n);  
    }  
  
    public synchronized void incrementValue(final int delta) {  
        holder = holder.incrementValue(delta);  
    }  
}
```

Locks and performances

- ▶ We know that declaring all methods of a class synchronized is a **simple solution to obtain thread-safety**.
- ▶ However, this technique can **compromise the performances** of the program.
- ▶ Extreme situations occur when only one thread at a time performs large synchronized operations, blocking all other threads.
- ▶ Mutexes can be very ineffective because threads may get suspended with consequent **context-switching**.

Uncontended synchronization

- ▶ When developing a concurrent application, the difference between **contended and uncontended synchronization** has to be taken into consideration.
 - ▶ Don't worry about uncontended synchronization. The Java compiler has **optimization support for it** (locks removal, escape analysis, ...).
- ▶ But even if the compiler is not able to optimize, **uncontended synchronization is not expensive** because threads are not frequently blocked at the lock.

Highly contended locks

- ▶ The main difficulties arise when there are **highly contended exclusive-locks**.
- ▶ **The serialization effect** caused by excessive synchronization reduces the concurrent execution potential. When a concurrent application runs serially, performances might be worse than for its mono-thread version.
- ▶ Two factors produce highly contended locks: **the number of requests** and the **duration of the lock**.
- ▶ In extreme cases, if many threads are waiting, computing resources might be underused even if there is a lot of work to do.

Poor concurrency

- ▶ We speak of **poor concurrency**, when the number of simultaneous invocations is limited by the program structure (instead of the available processing resources).

Locks and concurrency

- ▶ To increase concurrency, the size of the synchronized code blocks (or protected by other locking solutions) has to be reduced.
- ▶ **Balance between simplicity** (synchronized on the whole method) **and concurrency** (synchronized on small parts of code) has to be found.
- ▶ The acquisition of locks involves **overheads**, therefore splitting a synchronized code block into multiple ones is not always the advisable solution.
- ▶ **Compromises** have to be found on the different aspects of the application design.

Reducing the lock contention

There are 3 ways to reduce the lock contention:

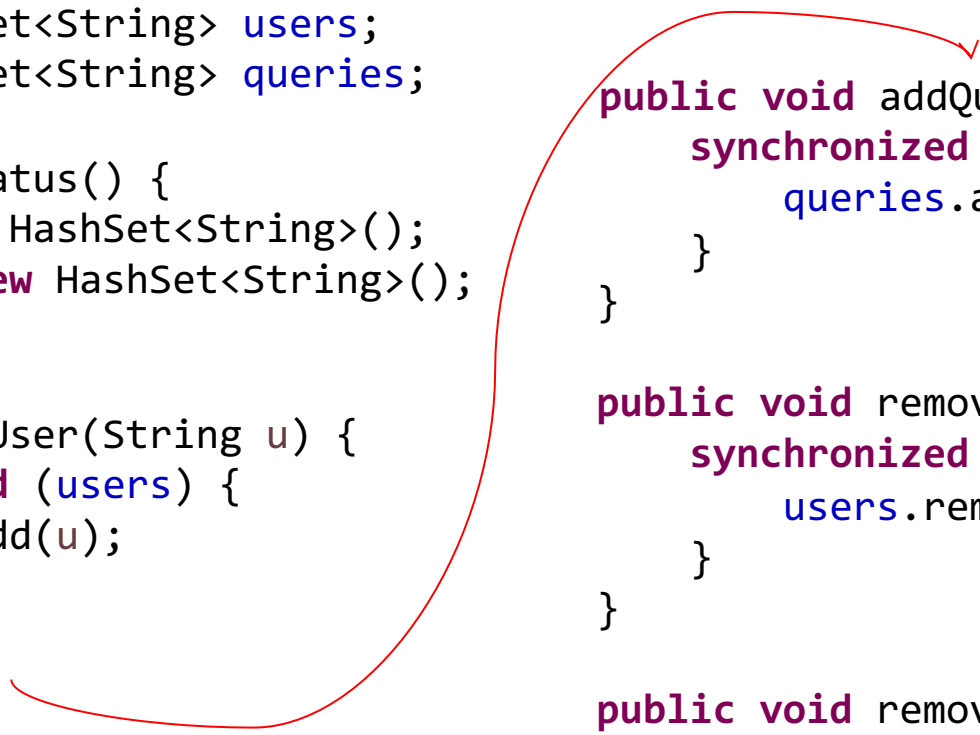
- ▶ Reduce the lock duration
 - ▶ less synchronized code
 - ▶ lock splitting
- ▶ Reduce the lock request frequency
- ▶ Replace exclusive locks with other synchronization approaches (immutability, volatile variables, atomic variables, ...).

Example: lock splitting

```
public class ServerStatus {  
    private final Set<String> users;  
    private final Set<String> queries;  
  
    public ServerStatus() {  
        users = new HashSet<String>();  
        queries = new HashSet<String>();  
    }  
  
    public synchronized void addUser(String u) {  
        users.add(u);  
    }  
  
    public synchronized void addQuery(String q) {  
        queries.add(q);  
    }  
  
    public synchronized void removeUser(String u) {  
        users.remove(u);  
    }  
  
    public synchronized void removeQuery(String q) {  
        queries.remove(q);  
    }  
}
```

Example: lock splitting

```
public class ServerStatus {  
    private final Set<String> users;  
    private final Set<String> queries;  
  
    public ServerStatus() {  
        users = new HashSet<String>();  
        queries = new HashSet<String>();  
    }  
  
    public void addUser(String u) {  
        synchronized (users) {  
            users.add(u);  
        }  
    }  
  
    public void addQuery(String q) {  
        synchronized (queries) {  
            queries.add(q);  
        }  
    }  
  
    public void removeUser(String u) {  
        synchronized (users) {  
            users.remove(u);  
        }  
    }  
  
    public void removeQuery(String q) {  
        synchronized (queries) {  
            queries.remove(q);  
        }  
    }  
}
```



- ▶ A program that takes advantage of lock splitting is a program that **has contended locks but uncontended data**.
- ▶ The granularity of the lock cannot be reduced if there are variables required for each operation **(hot fields)**.
- ▶ If hot fields are present, the only option is to modify the design of the program, by e.g. considering alternative solutions to the lock.

Alternatives to mutexes

- ▶ **Volatile variables**
 - ▶ for status variables (e.g. stop flags)
- ▶ **Immutable objects**
- ▶ **Various techniques for thread safety**
 - ▶ example: thread-confinement, stack-confinement, ...
- ▶ **ReadWrite locks**
 - ▶ acquire the lock only at the time of writing
 - ▶ very good if lot of reads and just a few writes are performed
- ▶ **Atomic Variables**
 - ▶ good for counters, management of dynamic data structures, ...
- ▶ **Concurrent collections and synchronizers**

Summary of topics

- ▶ Safe construction
- ▶ Publication of objects
- ▶ Safe publication
- ▶ Effectively immutable objects
- ▶ Publication and mutable/thread-safe objects
- ▶ Sharing of objects after publication
- ▶ Technique of the immutable holder
- ▶ Locks and performances
- ▶ Possible locks optimizations