

# Algorithms and Data Structures

## B Trees

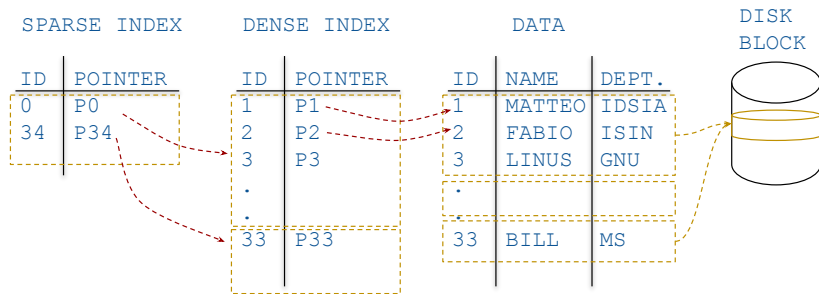
**Matteo Salani**  
matteo.salani@idsia.ch

# Indexing and Databases

The need of indexing:

- ▶ In databases, the amount of data to be stored is very high.
- ▶ Data cannot be stored in the main memory, hence it is commonly stored in the disk.
- ▶ Data access from the disk is much slower compared to the main memory access.
- ▶ Data is usually accessed in the form of blocks.

# Indexing Example



Goal of multi-level indexing:

- ▶ Reduce the number of blocks to process to access data
- ▶ Challenge: levels must grow and shrink dynamically

# B-Trees

- ▶ B-trees are balanced search trees designed to work with blocks of data.
- ▶ Many database systems use B-trees, or variants, to store information.
- ▶ B-tree nodes may have many children (from a few to thousands)
- ▶ The height of a B-tree can be considerably less than that of a BST





Figure: Example of a B-Tree

# B-Trees



1 node,  
1000 keys

1001 nodes,  
1,001,000 keys

1,002,001 nodes,  
1,002,001,000 keys

**Figure:** Example of a B-Tree with over one billion keys

# Definition of B-Trees

A B-tree  $T$  is a rooted tree (whose root is  $T.root$ ), let  $x$  be any node:

1.  $x.n$  is the number of keys stored in node  $x$
2. keys  $x.k_i$   $i \in \{1 \dots x.n\}$  are stored in non decreasing order  
( $x.k_1 \leq x.k_2 \leq \dots \leq x.k_{x.n}$ )
3. Each node has  $n + 1$  pointers  $x.c_i$   $i \in \{1 \dots x.n + 1\}$  (when  $x$  is a leaf, pointers have no meaning)
4. Keys separate the range of keys stored in subtrees  
 $k_{11} \dots k_{1n_1} \leq x.k_1 \leq k_{21} \dots k_{2n_2} \leq x.k_2 \leq \dots \leq x.k_{x.n} \leq k_{n1} \dots k_{nn_n}$
5. All leaves have the same depth




Figure: Example of a B-Tree

# Minimum degree or Order of B-Trees

Nodes have **lower** and **upper** bounds on the number of keys they can contain. Let  $t$  be the **minimum-degree** or **order** of a B-tree.

1. Every node (except the root) must have at least  $t - 1$  keys (thus at least  $t$  children).
2. Every node may contain at most  $2t - 1$  keys (thus at most  $2t$  children).

# Height of a B-Trees

**Theorem** If  $N \geq 1$ , the height  $h$  of a B-Tree  $T$  of degree  $t$  satisfies:

$$h \leq \log_t \frac{N+1}{2}$$

**Proof** The root has at least 1 key, the others nodes have at least  $t - 1$  keys.  $T$  has 2 nodes at level 1,  $2 \cdot t$  nodes at level 2,  $2 \cdot t^2$  nodes at level 3,  $2 \cdot t^{h-1}$  nodes at level  $h$ ,




Figure: Minimum number of nodes



# Height of a B-Trees

Number of nodes

$$1 + 2 + 2t + 2t^2 + \dots + 2t^{h-1} = 1 + \sum_{i=1}^h 2 \cdot t^{i-1}$$

Number of keys

$$\begin{aligned} N &\geq 1 + (t-1) \cdot \sum_{i=1}^h 2 \cdot t^{i-1} = 1 + 2 \cdot (t-1) \cdot \sum_{i=0}^{h-1} t^i \\ &= 1 + 2 \cdot (t-1) \left( \frac{1-t^h}{1-t} \right) = 1 + 2 \cdot (t-1) \left( \frac{t^h-1}{t-1} \right) \\ &= 1 + 2 \cdot (t^h-1) = 2 \cdot t^h - 1 \end{aligned}$$

$$N \geq 2 \cdot t^h - 1 \Rightarrow t^h \leq (N+1)/2 \Rightarrow \log_t t^h \leq \log_t (N+1)/2$$

$$h \leq \log_t (N+1)/2$$

# Searching in a B-Tree

```
1 def create(T):  
2     x = allocateNode()  
3     x.n = 0  
4     T.root = x
```

Listing 1: Create method

```
1 def search(x, key):  
2  
3     # Similar to BST  
4     i = 1  
5     while i <= x.n and key > x.key[i]:  
6         i = i + 1  
7     if i <= x.n and key == x.key[i]:  
8         return x, i  
9     else if isLeaf(x):  
10        return NULL  
11    else  
12        search(x.c[i], key)
```

Listing 2: Search method

# Inserting in a B-Tree


Inserting keys in a B-Tree is more complex than in BST.

- ▶ Keys are inserted in **existing** leaf nodes
- ▶ If the leaf node is **full** the node is **split** around the median key
- ▶ The median key moves to the parent node

```
1 def split(x, i):
2     z = allocateNode()
3     y = x.c[i]
4     z.n = t-1
5     for j = 1 to t-1:
6         z.key[j] = y.key[j+t]
7     if not isLeaf(y)
8         for j = 1 to t
9             z.c[j] = y.c[j+t]
10    y.n = t-1
11    for j=x.n+1 downto i+1
12        x.c[j+1] = x.c[j]
13    x.c[i+1] = z
14    for j=x.n downto i
15        x.key[j+1] = x.key[j]
16    x.key[i] = y.key[t]
17    x.n = x.n+1
```

Listing 3: Split child method

# Split Example



**Figure:** Split operation, median element "S" goes up one level


# Inserting in a B-Tree

We insert a key  $k$  into a B-tree  $T$  of height  $h$ . It requires  $O(h)$  block accesses and  $O(t \cdot h) = O(t \cdot \log_t n)$

```
1 def insert(T, k):
2     r = T.root
3     # Case split root
4     if r.n == 2t-1
5         s = allocateNode()
6         T.root = s
7         s.n = 0
8         s.c[1] = r
9         split(s,1)
10        insert_nonfull(s, k)
11    else
12        insert_nonfull(r, k)
13
```

Listing 4: Insert method

# B-Tree growing



**Figure:** A B-Tree grows from the top, splitting the root

# Inserting in a B-Tree

We insert a key  $k$  into a B-tree  $T$  of height  $h$ . It requires  $O(h)$  block accesses and  $O(t \cdot h) = O(t \cdot \log_t n)$

```
1 def insert_nonfull(x, k):
2     i = x.n
3     if isLeaf(x)
4         while i >= 1 and k < x.key[i]
5             x.key[i+1] = x.key[i]
6             i = i-1
7         x.key[i+1] = k
8         x.n = x.n+1
9     else
10        while i >= 1 and k < x.key[i]
11            i = i-1
12        i = i+1
13        if x.c[i].n == 2t-1
14            split(x,i)
15            if k > x.key[i]
16                i = i+1
17        insert_nonfull(x.c[i], k)
18
```

Listing 5: Insert non full method

# B-Tree example - Order 3

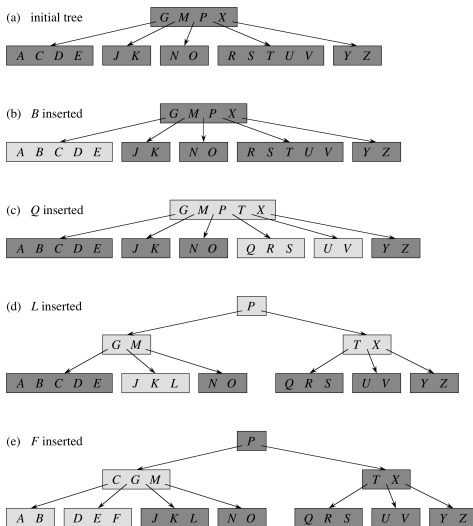


Figure: Sequence of insertion, B, Q, L, F



# Deleting in a B-Tree

- ▶ We can delete a key from any node
- ▶ We may have to rearrange the node's children.
- ▶ We must respect B-Tree properties

# Deleting in a B-Tree

Sketch of the algorithm, not a real pseudo code

```
1 def delete(x, k):
2     if k in x and isLeaf(x) # case 1
3         deleteElement(x, k)
4     if k in x and not isLeaf(x)
5         y = findPred(k,x) # Return the pred. node of k
6         if y.n >= t # case 2.a
7             # return the pred. key of k
8             k1 = findPredKey(k,y)
9             swap(k,k1)
10            delete(y, k)
11     else
12         z = findNext(k,x) # Return the next node of k
13         if z.n >= t # case 2.b
14             # return the next key of k
15             k1 = findNextKey(k,z)
16             swap(k,k1)
17             delete(z, k)
18         else # case 2.c
19             merge(y, k, z)
20             freeze(z)
21             delete(k, y)
22
```

# Deleting in a B-Tree


## Sketch of pseudo code (continued)

```
1  else # k is not in x
2      # find the subtree containig k
3      i = findChild(x, k)
4      if x.c[i].n == t-1
5          if i > 1 and x.c[i-1].n >= t # case 3.a
6              exchange(x.c[i-1], x)
7              delete(x.c[i], k)
8          else if i < x.n and x.c[i+1].n >= t # case 3.a
9              exchange(x.c[i+1], x)
10             delete(x.c[i], k)
11         else # case 3.b, index i must be checked
12             merge(x.c[i-1], x.key[i-1], x.c[i])
13             or merge(x.c[i], x.key[i], x.c[i+1])
14             delete(x, k)
15             # Take care of root node if it gets empty
16     else
17         delete(x.c[i], k)
18
```


Listing 7: Delete method

# B-Tree example - Order 3


(a) initial tree



(b) *F* deleted: case 1



(c) *M* deleted: case 2a



(d) *G* deleted: case 2c






Figure: Deletion of *F*, *M*, *G*

## B-Tree example - Order 3

(e) *D* deleted: case 3b



(e') tree shrinks  
in height



(f) *B* deleted: case 3a




Figure: Deletion of *D*, *B*